

```
In [1]: import pandas as pd
In [2]: import numpy as np
In [3]: df=pd.read_csv("C:/Users/daksh/OneDrive/Documents/Desktop/datasets/uber.csv")
In [4]: df.head()
Out[4]:
   Unnamed: 0      key  fare_amount  pickup_datetime  pickup_longitude  pickup_latitude  dropoff_longitude  dropoff_latitude  passenger_count
0    24238194  2015-05-07 19:52:06.0000003     7.5  2015-05-07 19:52:06 UTC   -73.999817    40.738354   -73.999512    40.723217           1
1    27835199  2009-07-17 20:04:56.0000002     7.7  2009-07-17 20:04:56 UTC   -73.994355    40.728225   -73.994710    40.750325           1
2    44984355  2009-08-24 21:45:00.0000001    12.9  2009-08-24 21:45:00 UTC   -74.005043    40.740770   -73.962565    40.772647           1
3    25894730  2009-06-26 08:22:21.0000001     5.3  2009-06-26 08:22:21 UTC   -73.976124    40.790844   -73.965316    40.803349           3
4    17610152  2014-08-28 17:47:00.000000188    16.0  2014-08-28 17:47:00 UTC   -73.925023    40.744085   -73.973082    40.761247           5

In [5]: df.shape
Out[5]: (200000, 9)
In [6]: df.isnull().sum()
Out[6]:
   Unnamed: 0      0
key            0
fare_amount    0
pickup_datetime 0
pickup_longitude 0
pickup_latitude 0
dropoff_longitude 1
dropoff_latitude 1
passenger_count 0
dtype: int64
In [7]: df = df.dropna(subset=['dropoff_latitude', 'dropoff_longitude'])
In [8]: df.isnull().sum()
Out[8]:
   Unnamed: 0      0
key            0
fare_amount    0
pickup_datetime 0
pickup_longitude 0
pickup_latitude 0
dropoff_longitude 0
dropoff_latitude 0
passenger_count 0
dtype: int64
In [9]: int(df['pickup_latitude'].isnull().sum())
df.dtypes
Out[9]:
   Unnamed: 0      int64
key            object
fare_amount    float64
pickup_datetime object
pickup_longitude float64
pickup_latitude float64
dropoff_longitude float64
dropoff_latitude float64
passenger_count int64
dtype: object
In [10]: df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
df.dtypes
Out[10]:
   Unnamed: 0      int64
key            object
fare_amount    float64
pickup_datetime datetime64[ns, UTC]
pickup_longitude float64
pickup_latitude float64
dropoff_longitude float64
dropoff_latitude float64
passenger_count int64
dtype: object
In [11]: # Remove invalid latitude and longitude values
df = df[
    (df['pickup_latitude'].between(-90, 90)) &
    (df['dropoff_latitude'].between(-90, 90)) &
    (df['pickup_longitude'].between(-180, 180)) &
    (df['dropoff_longitude'].between(-180, 180))
]

In [12]: from geopy.distance import great_circle
# Compute haversine distance
def calculate_distance(row):
    pickup = (row['pickup_latitude'], row['pickup_longitude'])
    dropoff = (row['dropoff_latitude'], row['dropoff_longitude'])
    return great_circle(pickup, dropoff).km

df['distance_km'] = df.apply(calculate_distance, axis=1)

# Remove extreme values
df = df[(df['distance_km'] > 0) & (df['distance_km'] < 50)]
df = df[df['fare_amount'] < 100]

#geopy is a Python library that makes it easy to calculate distances, latitudes, and longitudes between two points on the surface of the Earth
#The great circle distance formula provides an accurate calculation of the distance between two points on the Earth's surface.

In [13]: df.shape
Out[13]: (193797, 10)
In [14]: print(df.columns)
Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
       'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
       'dropoff_latitude', 'passenger_count', 'distance_km'],
       dtype='object')
In [15]: from geopy.distance import great_circle
def calculate_distance(row):
    try:
        pickup = (row['pickup_latitude'], row['pickup_longitude'])
        dropoff = (row['dropoff_latitude'], row['dropoff_longitude'])
        return great_circle(pickup, dropoff).km
    except:
        return None

df['distance_km'] = df.apply(calculate_distance, axis=1)
df = df.dropna(subset=['distance_km'])

In [21]: import seaborn as sns
import matplotlib.pyplot as plt

corr = df[['fare_amount', 'distance_km']].corr()
sns.heatmap(corr, annot=True)
plt.title("Correlation Matrix")
plt.show()



The Correlation Matrix heatmap shows the correlation between fare_amount and distance_km. The diagonal elements are 1. The off-diagonal elements show a strong positive correlation, with a value of approximately 0.89.


In [22]: # Step 4: Train Linear Regression & Random Forest Models
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor

In [27]: # Select features and target
X = df[['distance_km']]
y = df[['fare_amount']]

In [28]: # Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [29]: # Linear Regression
lr = LinearRegression()
lr.fit(X_train, y_train)
lr_preds = lr.predict(X_test)

In [30]: # Random Forest Regression
rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
rf_preds = rf.predict(X_test)

C:\Users\daksh\AppData\Local\Programs\Python\Python314\Lib\site-packages\sklearn\base.py:1365: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,) for example using ravel().
return fit_method(estimator, *args, **kwargs)

In [31]: # Step 5: Evaluate the Models
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

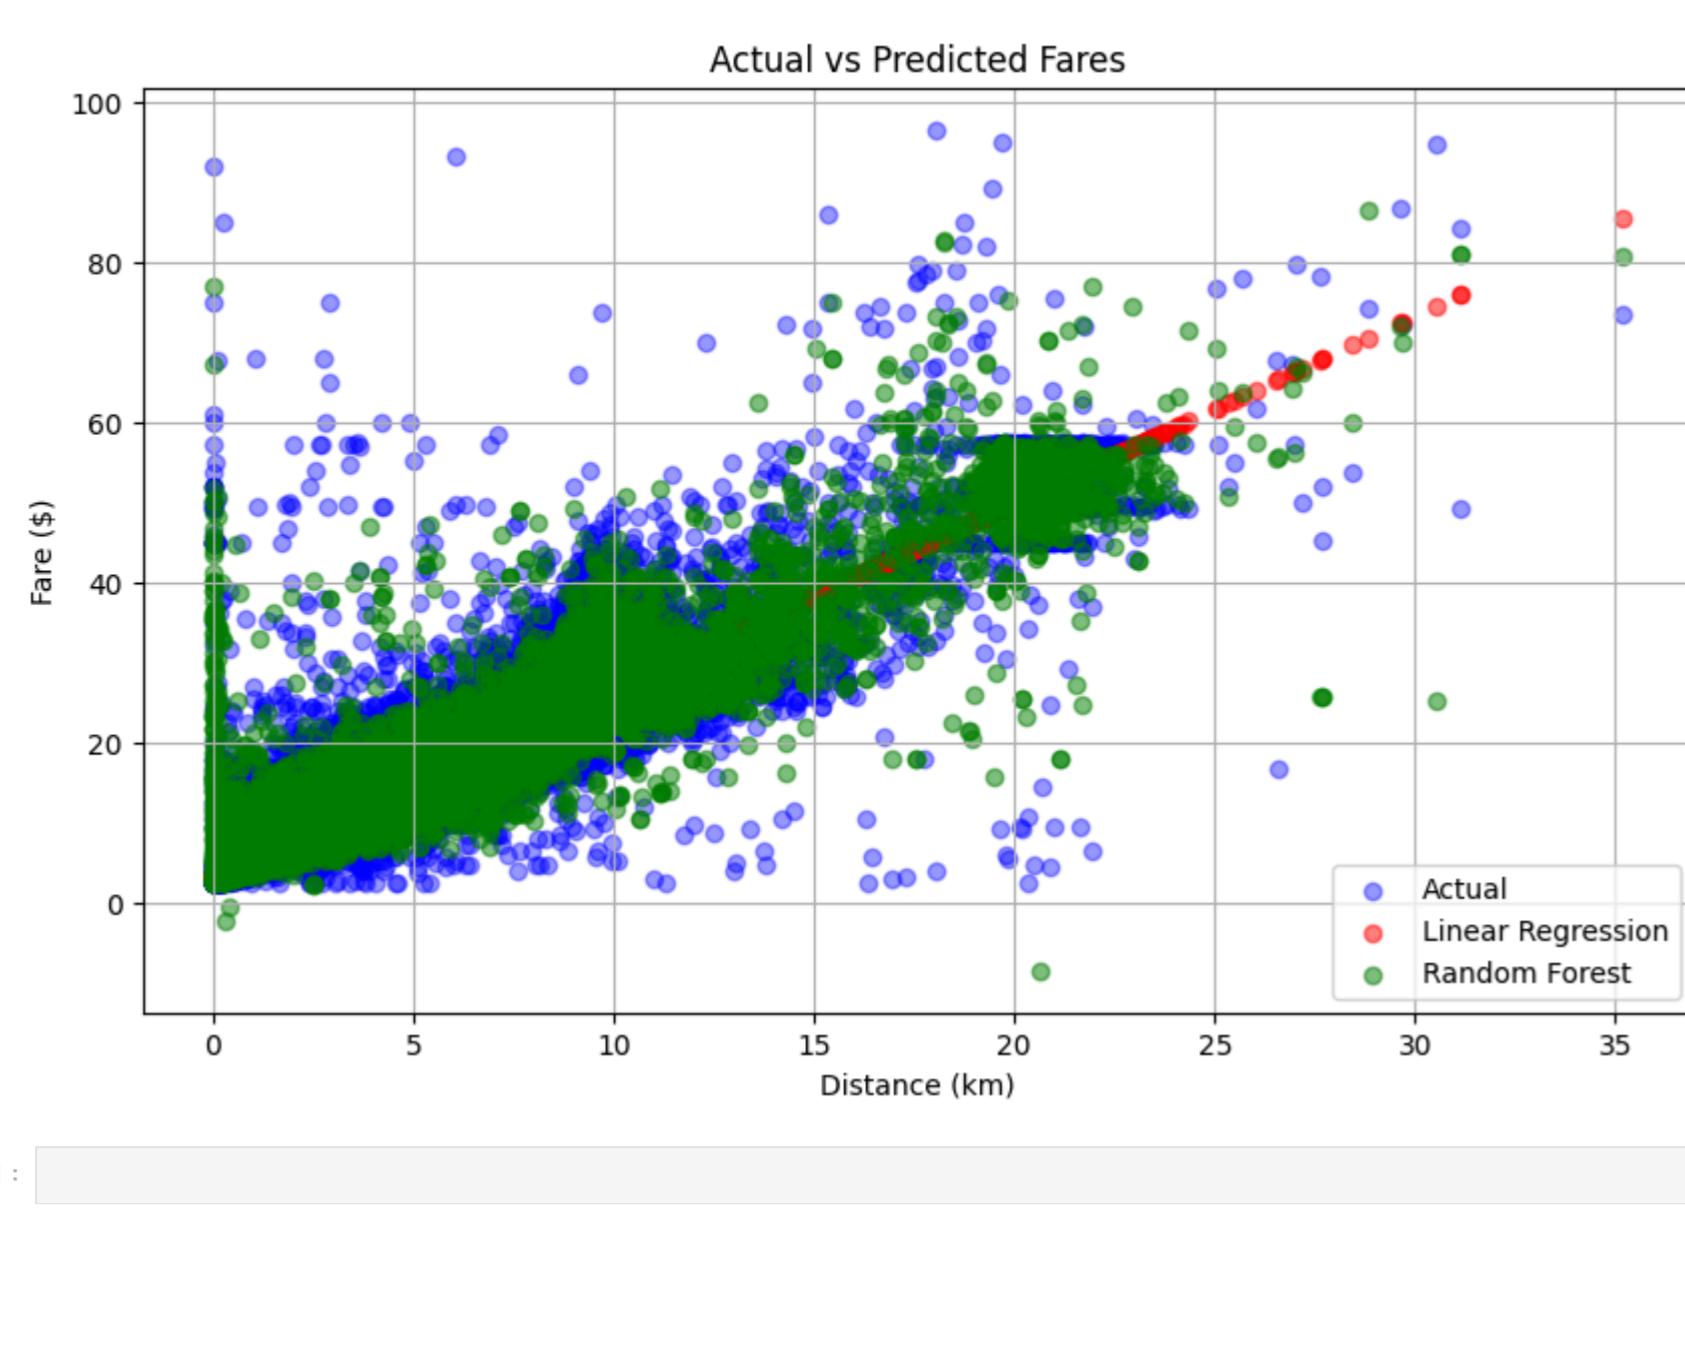
def evaluate(y_true, y_pred, model_name):
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    r2 = r2_score(y_true, y_pred)
    print(f"{model_name} Results:")
    print(f"RMSE: {rmse:.2f}")
    print(f"R^2 Score: {r2:.2f}\n")

evaluate(y_test, lr_preds, "Linear Regression")
evaluate(y_test, rf_preds, "Random Forest Regression")

Linear Regression Results:
RMSE: 4.20
R^2 Score: 0.80

Random Forest Regression Results:
RMSE: 4.99
R^2 Score: 0.71

In [33]: # Visual Comparison
plt.figure(figsize=(10, 6))
plt.scatter(X_test, y_test, label="Actual", color='blue', alpha=0.4)
plt.scatter(X_test, lr_preds, label="Linear Regression", color='red', alpha=0.5)
plt.scatter(X_test, rf_preds, label="Random Forest", color='green', alpha=0.5)
plt.xlabel("Distance (km)")
plt.ylabel("Fare ($)")
plt.title("Actual vs Predicted Fares")
plt.legend()
plt.grid(True)
plt.show()
```



In [ ]: